

CS 260: Foundations of Data Science

Prof. Thao Nguyen

Fall 2024



Haverford
COLLEGE

Admin

- **Lab 2** grades & feedback will be posted on Wednesday
- **Lab 3** due tonight
- **Lab 4** posted, due next Monday at midnight
- **Lecture Notes**

Peer Tutoring

- **Student tutors** (Fejiro Anigbro, Darshan Mehta)
- **Flexible hours**
- **Free!**



TECH TALKS

2024

OCTOBER 7, 8 & 9TH | 6-8PM EST

***Sign up for a 30 minute virtual informational interview
with a Tri-Co alum to gain tech career insights!***

Alumni will represent various tech roles including software engineering and development, data science, tech consulting, product management and biotech.

OCT 7	OCT 8	OCT 9
Accenture	Bristol Myers Squibb	The Walt Disney Company
•	•	•
FERMAT Commerce	Community.com	Fresh Tracks Insights
•	•	•
Grubhub	C3 Presents (Live Nation)	Meta
	•	•
	Opower (Oracle)	Grubhub

TRI-COLLEGE RECRUITING CONSORTIUM

HAVERFORD

BRYN MAWR

SWARTHMORE

Outline for today

- Recap SGD (stochastic gradient descent)
- Introduction to classification
 - Decision tree models
 - Probabilistic interpretation
- Evaluation Metrics
 - Confusion matrices
 - Precision and recall
 - ROC curves

Outline for today

- Recap SGD (stochastic gradient descent)
- Introduction to classification
 - Decision tree models
 - Probabilistic interpretation
- Evaluation Metrics
 - Confusion matrices
 - Precision and recall
 - ROC curves

Stochastic Gradient Descent for Linear Regression

Key Idea: take the derivative of **one datapoint** at a time and use that to update w

$$\nabla J = \begin{bmatrix} \frac{\partial J}{\partial w_0} \\ \vdots \\ \frac{\partial J}{\partial w_p} \end{bmatrix}$$

Handout 6
142

derivative
wrt
example

$$J(\vec{w}) = \frac{1}{2} \sum_{i=1}^n (\underbrace{\vec{w} \cdot \vec{x}_i}_{\text{pred}} - \underbrace{y_i}_{\text{truth}})^2$$

derivative is very
large + unstable

$$\nabla J(\vec{w})_{\vec{x}_i} = (\vec{w} \cdot \vec{x}_i - y_i) \vec{x}_i$$

datapoint scalar vector

$$\begin{bmatrix} x_{i1} \\ x_{i2} \\ \vdots \\ x_{ip} \end{bmatrix}$$

Stochastic Gradient Descent for Linear Regression

SGD

Start with $\vec{w} = \vec{0}$ (vector of zeros)

while (epoch) iteration t :

for $i = 1, 2, \dots, n$: (shuffle)

$$\vec{w} \leftarrow \vec{w} - \alpha (\vec{w} \cdot \vec{x}_i - y_i) \vec{x}_i$$

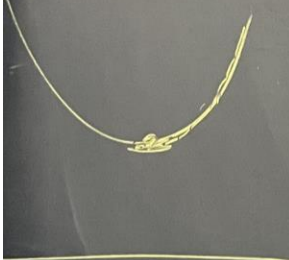
↑
Step size derivative

check convergence

if $|J(\vec{w}^t) - J(\vec{w}^{t+1})| < \epsilon$ $\leftarrow \epsilon$ is very small

\Rightarrow Stop!

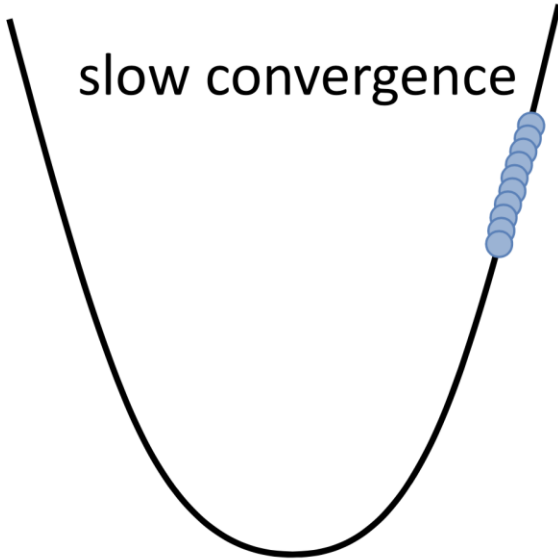
not for Lab 3



Choosing the step size alpha

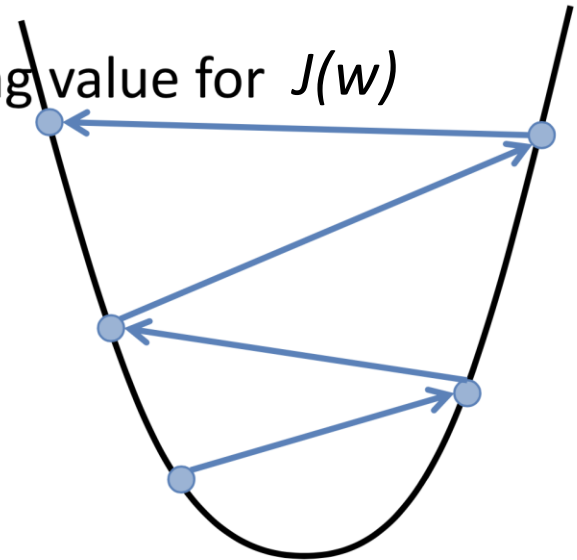
α too small

slow convergence



α too large

increasing value for $J(w)$



- may overshoot minimum
- may fail to converge (may even diverge)

Pros and Cons

(Analytic Solution)

Gradient Descent

- requires multiple iterations
- need to choose α
- works well when p is large
- can support online learning

Normal Equations

- non-iterative
- no need for α
- slow if p is large
 - matrix inversion is $O(p^3)$

Outline for today

- Recap SGD (stochastic gradient descent)
- Introduction to classification
 - Decision tree models
 - Probabilistic interpretation
- Evaluation Metrics
 - Confusion matrices
 - Precision and recall
 - ROC curves

Binary classification examples

- Transactions that indicate credit card fraud
- Accounts that are bots
- Detecting which scans show tumors
- Prenatal test for Down's Syndrome
- Finding genes under natural selection
- Regions of the environment that contains the object the robot is searching for

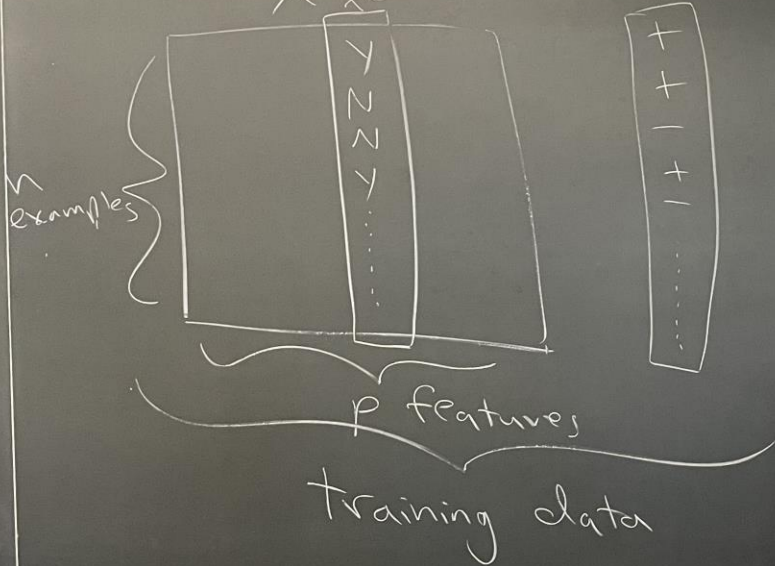
In all these examples, we are trying to find unusual items ("needle in a haystack") -- we call these *positives*

Introduction to Classification

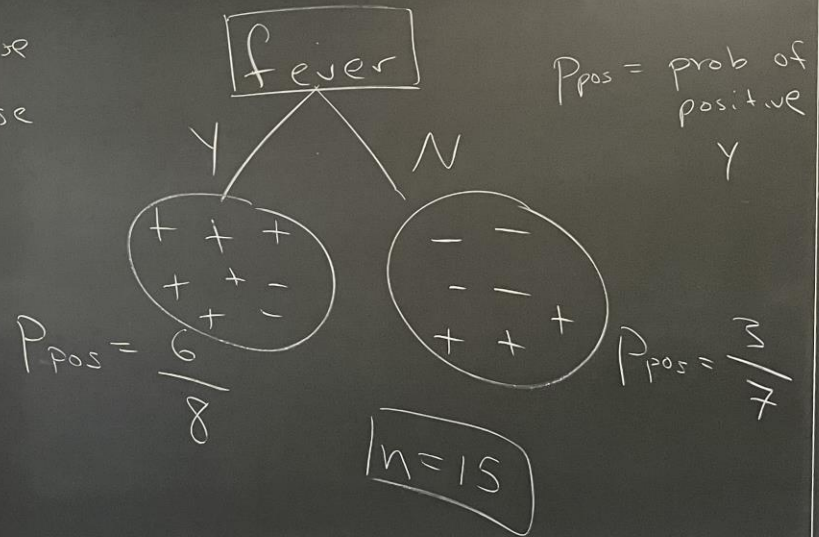
Classification

X fever

Y (disease)



model: decision tree with a single feature ("stump")



Introduction to Classification

new idea : use probabilities
to classify test examples

$$\vec{X}_{\text{test}} = \begin{bmatrix} \dots & \text{fever} & N & \dots \end{bmatrix}^T$$

threshold 0.5 \Rightarrow

$$\hat{y}_{\text{test}} = \ominus$$

no
disease

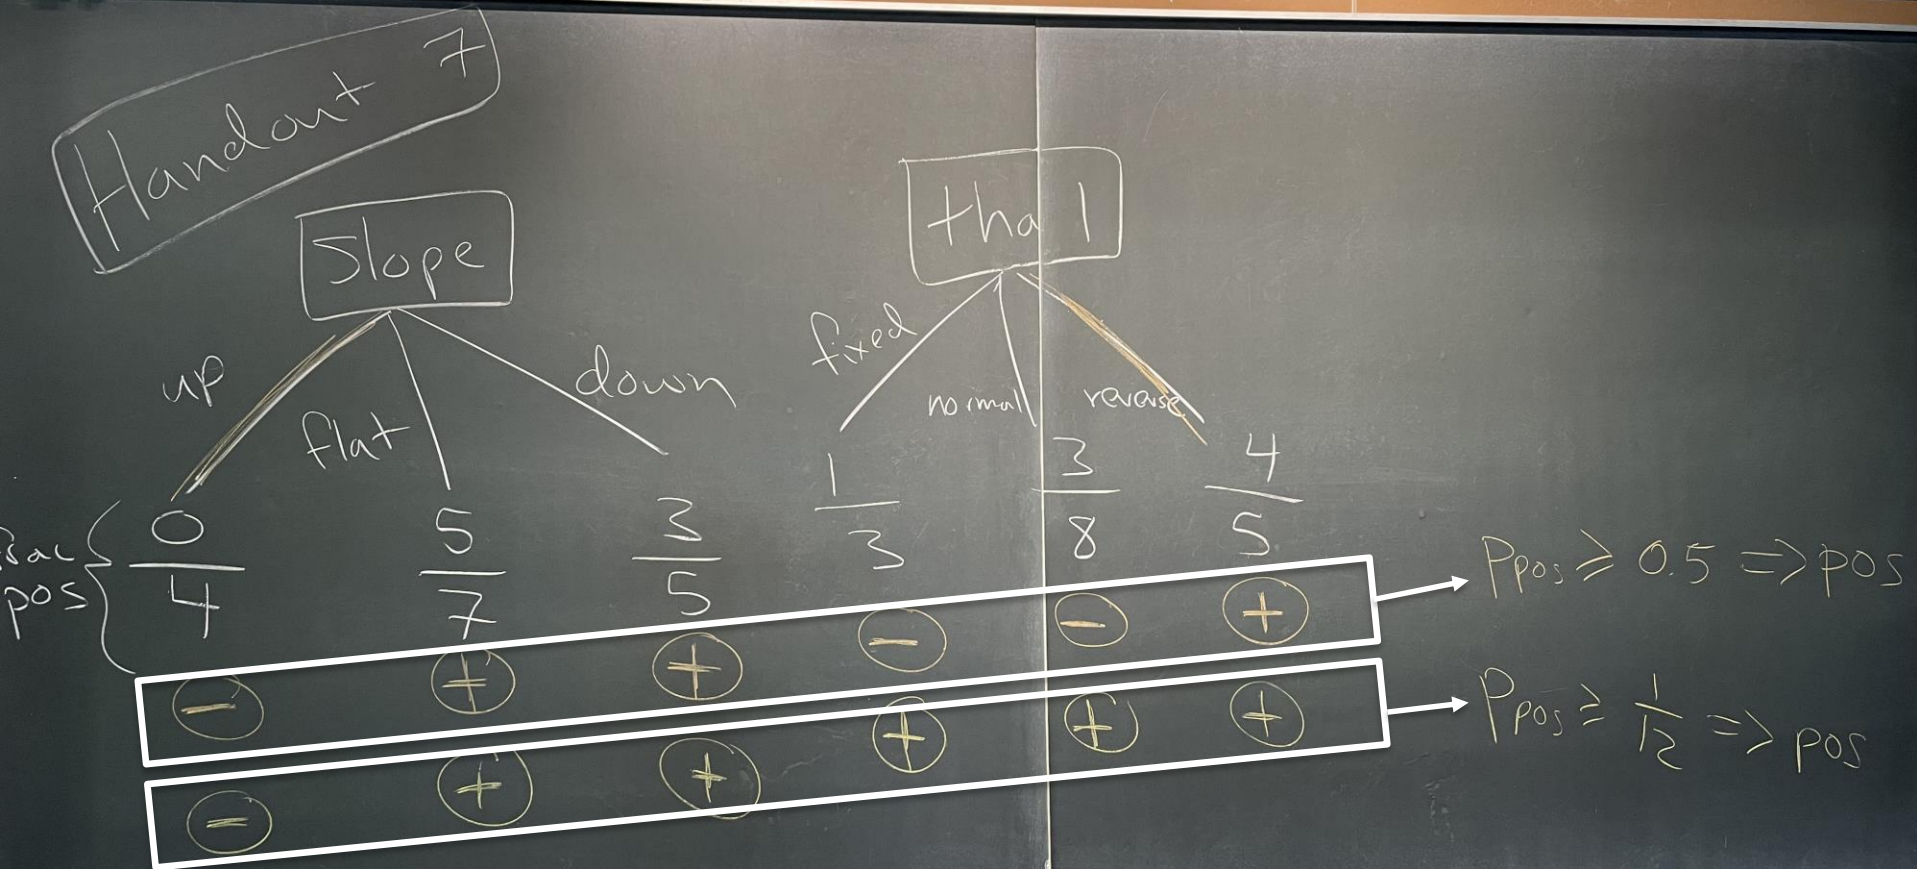
threshold 0.25 \Rightarrow

$$\hat{y}_{\text{test}} = \oplus$$

disease

$P_{\text{pos}} \geq \text{threshold} \Rightarrow \text{classify } \oplus$

Handout 7



Outline for today

- Recap SGD (stochastic gradient descent)
- Introduction to classification
 - Decision tree models
 - Probabilistic interpretation
- **Evaluation Metrics**
 - Confusion matrices
 - Precision and recall
 - ROC curves

Goals of Evaluation

- Think about what metrics are important for the problem at hand
- Compare different methods or models on the same problem
- Common set of tools that other researchers/users can understand

Training and Testing

(high-level idea)

- **Separate** data into “**train**” and “**test**”
 - n = num training examples
 - m = num testing examples
- **Fit** (create) the model using **training data**
 - e.g. sea_ice_1979-2012.csv
- **Evaluate** the model using **testing data**
 - e.g. sea_ice_2013-2020.csv

$$\frac{65+13}{100} = 78\%$$

pred

	-	+
truth	65	15
	7	13

Accuracy =

Note: all the same model,
different thresholds!

$$= \frac{1}{m} \sum_{i=1}^m \mathbb{1}(\hat{y}_i = y_i)$$

test data

$m = 100$

thresh
= 0.5

50	30
1	19

thresh
0.25

acc = 69%

80 negatives
20 positives

76	4
11	9

thresh
0.75

Confusion Matrices

		Predicted class	
		Negative	Positive
True class	Negative	True negative (TN)	False positive (FP)
	Positive	False negative (FN)	True positive (TP)

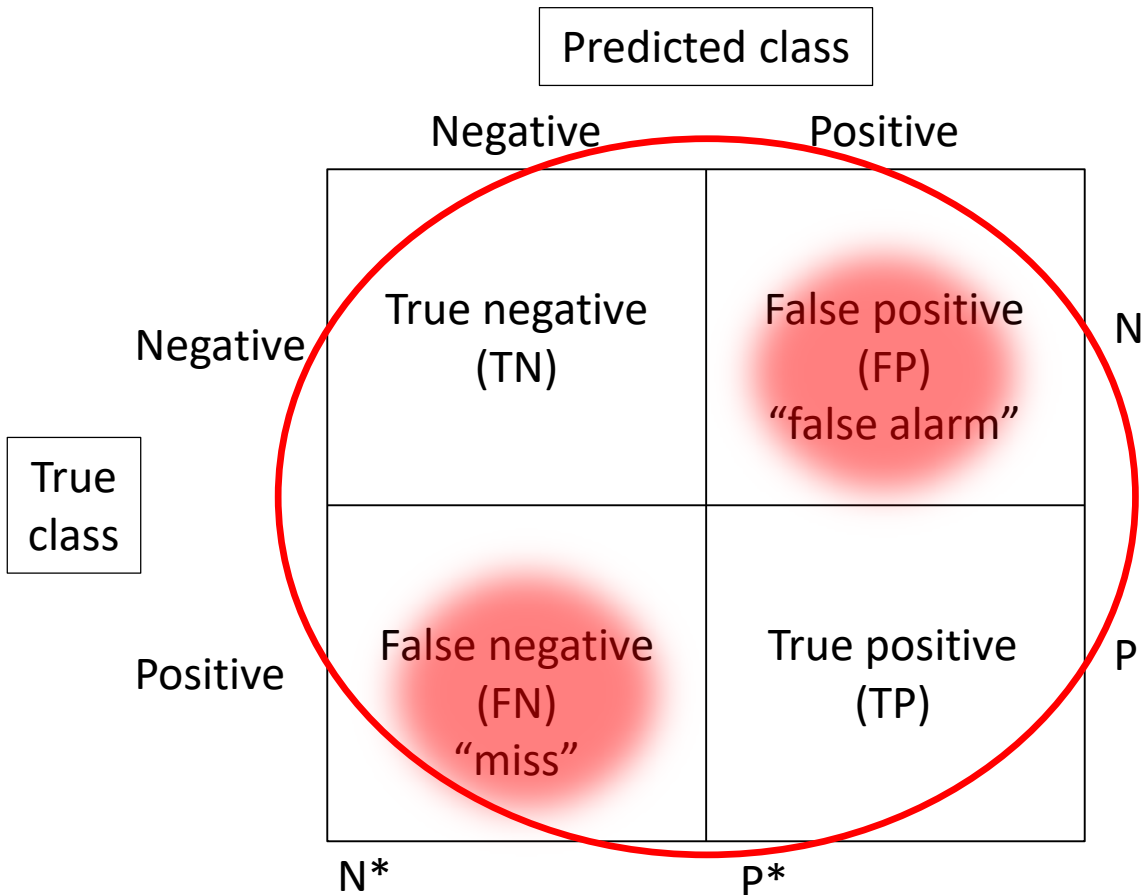
Confusion Matrices

		Predicted class		
		Negative	Positive	
True class	Negative	True negative (TN)	False positive (FP) “false alarm”	N (total number of true negatives)
	Positive	False negative (FN) “miss”	True positive (TP)	P (total number of true positives)
		N* (what we said was negative)	P* (what we said was positive “flagged”)	

Confusion Matrices

		Predicted class		
		Negative	Positive	
True class	Negative	True negative (TN) ✓	False positive (FP) “false alarm” ✗	N
	Positive	False negative (FN) “miss” ✗	True positive (TP) ✓	P
		N*	p*	

Confusion Matrices

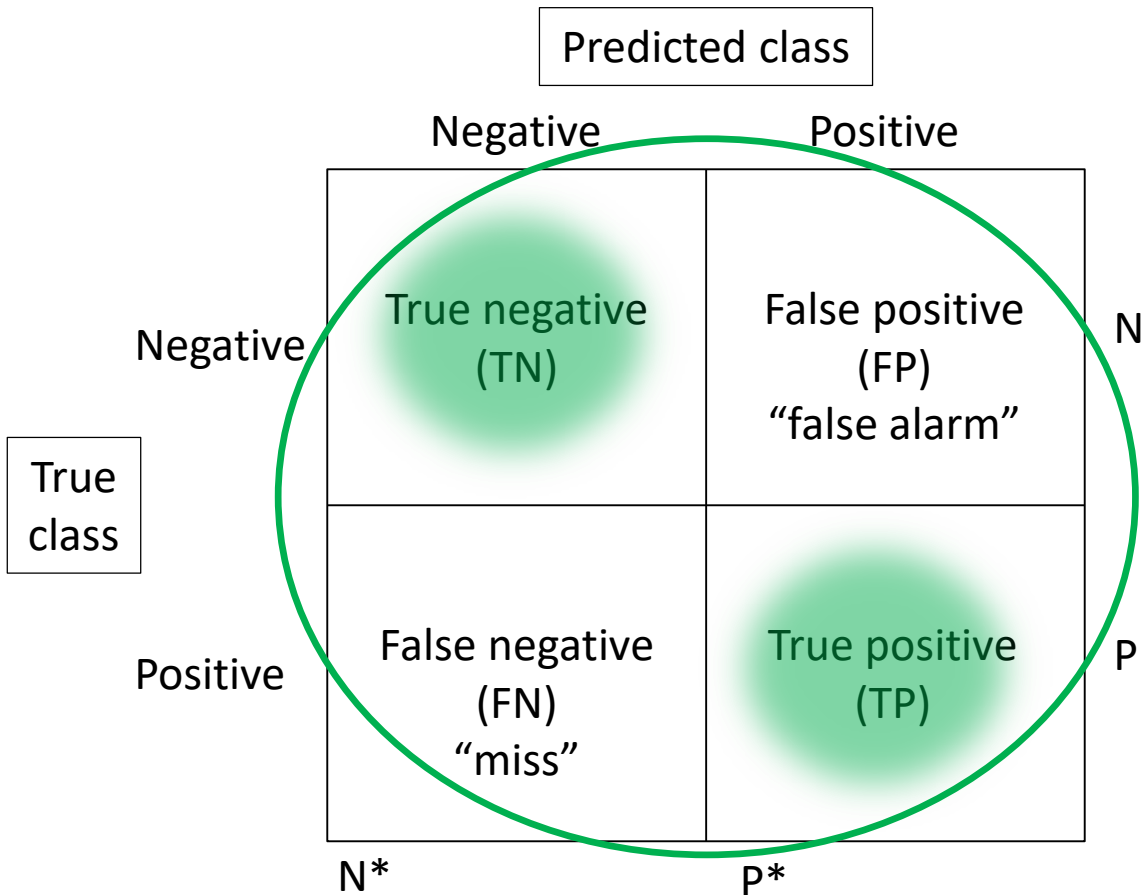


Error:

$$(FN+FP)/(TN+FP+FN+TP)$$

$$= (FN+FP)/(N+P)$$

Confusion Matrices



Accuracy = 1-Error:

$$(TN+TP)/(TN+FP+FN+TP)$$

$$= (TN+TP)/(N+P)$$

Confusion Matrices

		Predicted class		
		Negative	Positive	
True class	Negative	True negative (TN)	False positive (FP) “false alarm”	N
	Positive	False negative (FN) “miss”	True positive (TP)	P
		N*	p*	

Precision:

$$TP/(FP+TP) = TP/P^*$$

Confusion Matrices

		Predicted class		
		Negative	Positive	
True class	Negative	True negative (TN)	False positive (FP) “false alarm”	N
	Positive	False negative (FN) “miss”	True positive (TP)	P
		N*	p*	

Recall
(True Positive Rate):

$$TP/(FN+TP) = TP/P$$

Confusion Matrices

		Predicted class		
		Negative	Positive	
True class	Negative	True negative (TN)	False positive (FP) "false alarm"	N
	Positive	False negative (FN) "miss"	True positive (TP)	P
		N*	p*	

False Positive Rate:

$$FP/(TN+FP) = FP/N$$

Precision and Recall

- Precision: of all the “flagged” examples, which ones are actually relevant (i.e. positive)?

(Purity)

- Recall: of all the relevant results, which ones did I actually return?

(Completeness)