

CS 260: Foundations of Data Science

Prof. Thao Nguyen

Fall 2024



HAVERFORD
COLLEGE

Admin

- **Lab 2** grades & feedback will be posted on Wednesday
- **Lab 3** due tonight
- **Lab 4** posted, due next Monday at midnight
- **Lecture Notes**

Peer Tutoring

- **Student tutors** (Fejiro Anigbro, Darshan Mehta)
- **Flexible hours**
- **Free!**

TECH TALKS 2024

OCTOBER 7, 8 & 9TH | 6-8PM EST

***Sign up for a 30 minute virtual informational interview
with a Tri-Co alum to gain tech career insights!***

Alumni will represent various tech roles including software engineering and development, data science, tech consulting, product management and biotech.

| OCT 7 | OCT 8 | OCT 9 |
|-----------------|---------------------------|-------------------------|
| Accenture | Bristol Myers Squibb | The Walt Disney Company |
| • | • | • |
| FERMAT Commerce | Community.com | Fresh Tracks Insights |
| • | • | • |
| Grubhub | C3 Presents (Live Nation) | Meta |
| | • | • |
| | Opower (Oracle) | Grubhub |

TRI-COLLEGE RECRUITING CONSORTIUM

HAVERFORD

BRYN MAWR

SWARTHMORE

Outline for today

- Recap SGD (stochastic gradient descent)
- Introduction to classification
 - Decision tree models
 - Probabilistic interpretation
- Evaluation Metrics
 - Confusion matrices
 - Precision and recall
 - ROC curves

Outline for today

- Recap SGD (stochastic gradient descent)
- Introduction to classification
 - Decision tree models
 - Probabilistic interpretation
- Evaluation Metrics
 - Confusion matrices
 - Precision and recall
 - ROC curves

Stochastic Gradient Descent for Linear Regression

Key Idea: take the derivative of **one datapoint** at a time and use that to update w

$$\nabla J = \begin{bmatrix} \frac{\partial J}{\partial w_0} \\ \vdots \\ \frac{\partial J}{\partial w_p} \end{bmatrix}$$

Handout 6
142

derivative
wrt
example

$$J(\vec{w}) = \frac{1}{2} \sum_{i=1}^n (\underbrace{\vec{w} \cdot \vec{x}_i}_{\text{pred}} - \underbrace{y_i}_{\text{truth}})^2$$

derivative is very
large + unstable

$$\nabla J(\vec{w})_{\vec{x}_i} = (\vec{w} \cdot \vec{x}_i - y_i) \vec{x}_i$$

datapoint scalar vector

$$\begin{bmatrix} x_{i1} \\ x_{i2} \\ \vdots \\ x_{ip} \end{bmatrix}$$

Stochastic Gradient Descent for Linear Regression

SGD

Start with $\vec{w} = \vec{0}$ (vector of zeros)

while (epoch) iteration t :

for $i = 1, 2, \dots, n$: (shuffle)

$$\vec{w} \leftarrow \vec{w} - \alpha (\vec{w} \cdot \vec{x}_i - y_i) \vec{x}_i$$

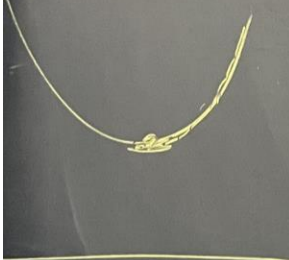
↑
Step size derivative

check convergence

if $|J(\vec{w}^t) - J(\vec{w}^{t+1})| < \epsilon$ $\leftarrow \epsilon$ is very small

\Rightarrow Stop!

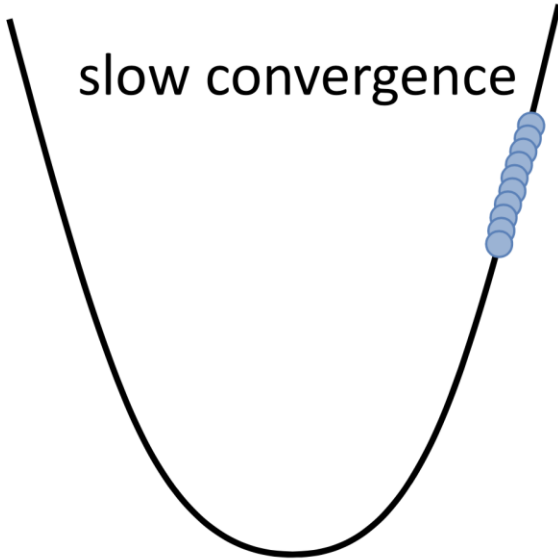
not for Lab 3



Choosing the step size alpha


α too small

slow convergence



α too large

increasing value for $J(w)$



- may overshoot minimum
- may fail to converge (may even diverge)

Pros and Cons

(Analytic Solution)

Gradient Descent

- requires multiple iterations
- need to choose α
- works well when p is large
- can support online learning

Normal Equations

- non-iterative
- no need for α
- slow if p is large
 - matrix inversion is $O(p^3)$

Outline for today

- Recap SGD (stochastic gradient descent)
- Introduction to classification
 - Decision tree models
 - Probabilistic interpretation
- Evaluation Metrics
 - Confusion matrices
 - Precision and recall
 - ROC curves

Binary classification examples

- Transactions that indicate credit card fraud
- Accounts that are bots
- Detecting which scans show tumors
- Prenatal test for Down's Syndrome
- Finding genes under natural selection
- Finding regions of the genome with high recombination rate (“hotspots”)

In all these examples, we are trying to find unusual items (“needle in a haystack”) -- we call these *positives*



Introduction to Classification

Classification

X fever

Y (disease)

n examples



+ \Rightarrow disease
- \Rightarrow no disease

p features

training data

model: decision tree with a single feature ("stump")

fever

Y

N



P_{pos} = prob of positive Y

$$P_{pos} = \frac{6}{8}$$

$$P_{pos} = \frac{3}{7}$$

n=15

Introduction to Classification

new idea : use probabilities
to classify test examples

$$\vec{X}_{\text{test}} = \begin{bmatrix} \dots & \text{fever} & N & \dots \end{bmatrix}^T$$

threshold 0.5 \Rightarrow

$$\hat{y}_{\text{test}} = \ominus$$

no
disease

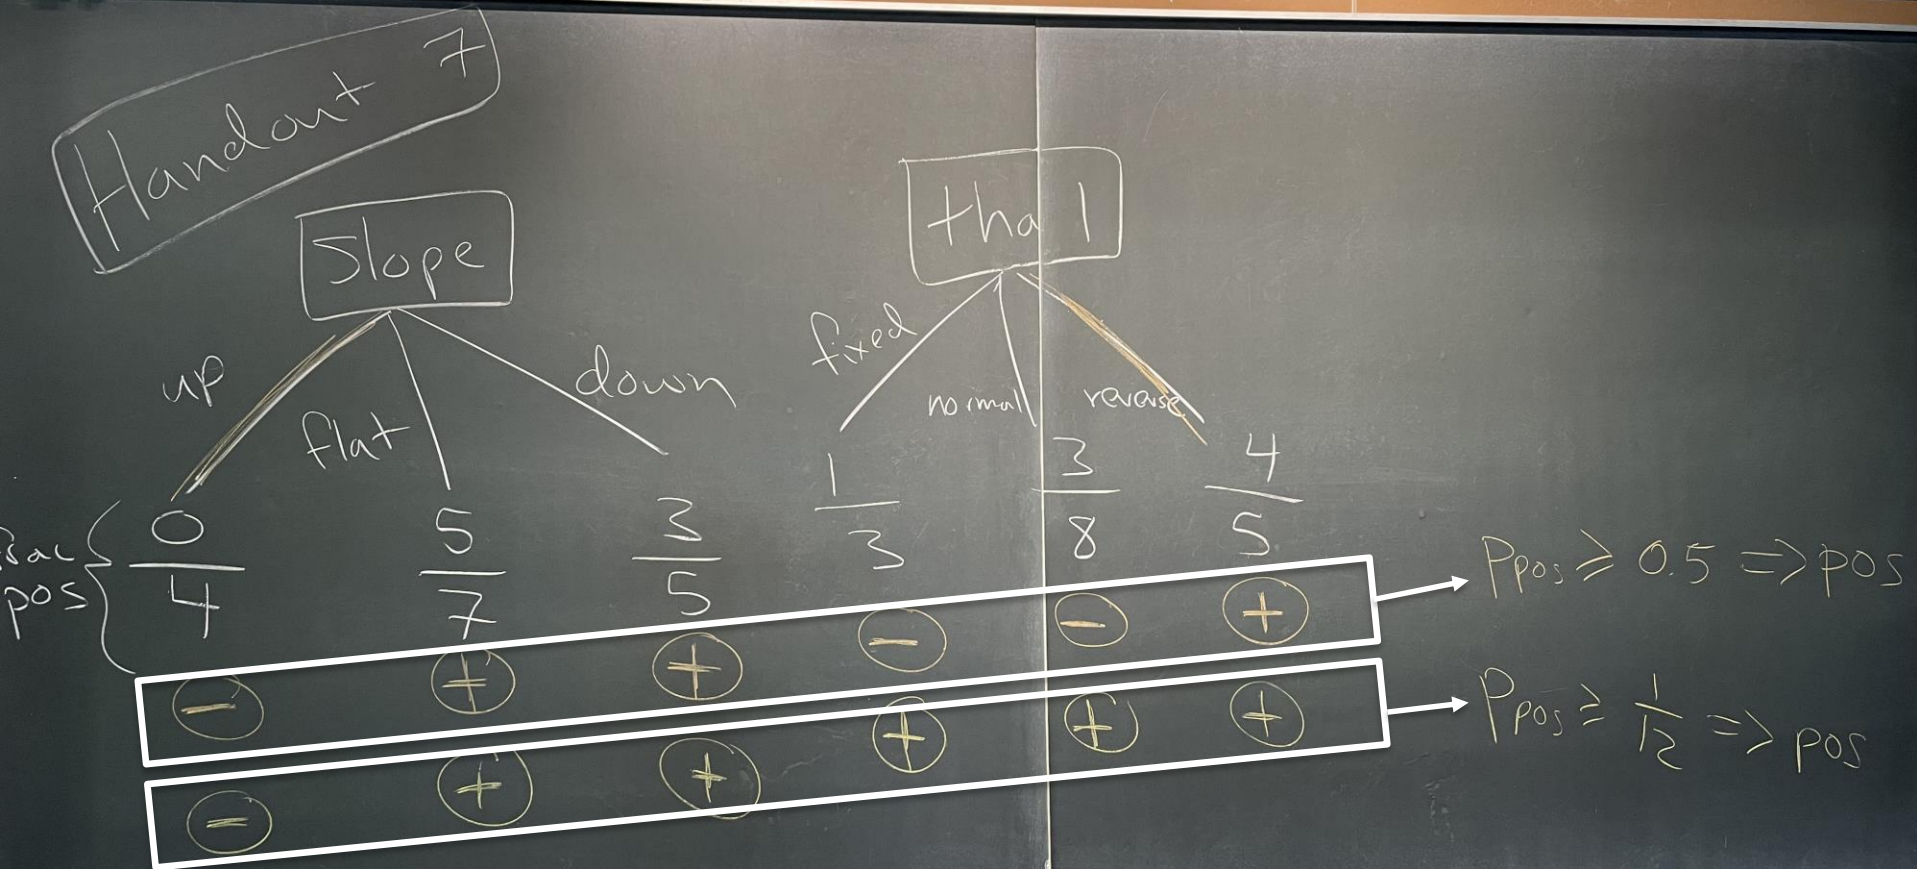
threshold 0.25 \Rightarrow

$$\hat{y}_{\text{test}} = \oplus$$

disease

$$P_{\text{pos}} \geq \text{threshold} \Rightarrow \text{classify } \oplus$$

Handout 7



Outline for today

- Recap SGD (stochastic gradient descent)
- Introduction to classification
 - Decision tree models
 - Probabilistic interpretation
- **Evaluation Metrics**
 - Confusion matrices
 - Precision and recall
 - ROC curves

Goals of Evaluation

- Think about what metrics are important for the problem at hand
- Compare different methods or models on the same problem
- Common set of tools that other researchers/users can understand

Training and Testing

(high-level idea)

- **Separate** data into “**train**” and “**test**”
 - n = num training examples
 - m = num testing examples
- **Fit** (create) the model using **training data**
 - e.g. sea_ice_1979-2012.csv
- **Evaluate** the model using **testing data**
 - e.g. sea_ice_2013-2020.csv

$$\frac{65+13}{100} = 78\%$$

Pred

| | - | + |
|-------|----|----|
| truth | 65 | 15 |
| | 7 | 13 |

Accuracy =

correct

$$= \frac{1}{m} \sum_{i=1}^m \mathbb{1}(\hat{y}_i = y_i)$$

Note: all the same model,
different thresholds!

test data

$m = 100$

Thresh
= 0.5

| | |
|----|----|
| 50 | 30 |
| 1 | 19 |

Thresh
0.25

acc = 69%

80 negatives
20 positives

| | |
|----|---|
| 76 | 4 |
| 11 | 9 |

Thresh
0.75

Confusion Matrices

| | | Predicted class | |
|------------|----------|---------------------|---------------------|
| | | Negative | Positive |
| True class | Negative | True negative (TN) | False positive (FP) |
| | Positive | False negative (FN) | True positive (TP) |